

x86 Assembly

Adapted from CS 161 Spring 2022 - Lecture 3

Next: x86 Assembly and Call Stack

- Part CS 61C review
 - How do computers represent numbers as bits and bytes?
 - How do computers interpret and run the programs we write?
 - How do computers organize segments of memory?
- Part new content
 - How does x86 assembly work?
 - How do you call a function in x86?

Number Representation

Textbook Chapter 2.1

Units of Measurement

- In computers, all data is represented as bits
 - **Bit**: a binary digit, 0 or 1
- Names for groups of bits
 - 4 bits = 1 **nibble**
 - 8 bits = 1 **byte**
- How many bits/nibbles/bytes in `0b 1000 1000 1000 1000`?
 - 16 bits, or 4 nibbles, or 2 bytes

Hexadecimal

- 4 bits can be represented as 1 hexadecimal digit (base 16)
- How would you write `0b11000110` in hex?
 - `0xC6`
 - Note: For clarity, we add `0b` in front of bits and `0x` in front of hex

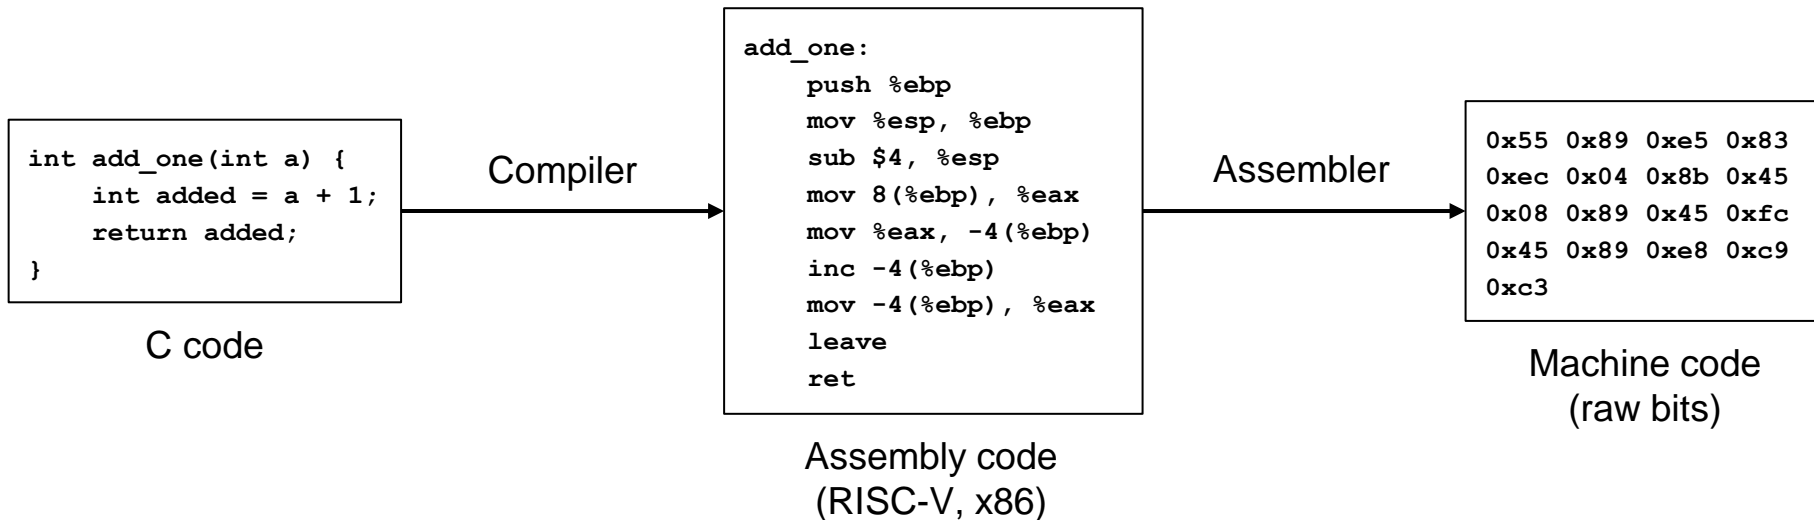
Binary	Hexadecimal
0000	0
0001	1
0010	2
0011	3
0100	4
0101	5
0110	6
0111	7

Binary	Hexadecimal
1000	8
1001	9
1010	A
1011	B
1100	C
1101	D
1110	E
1111	F

Running C Programs

Textbook Chapter 2.2

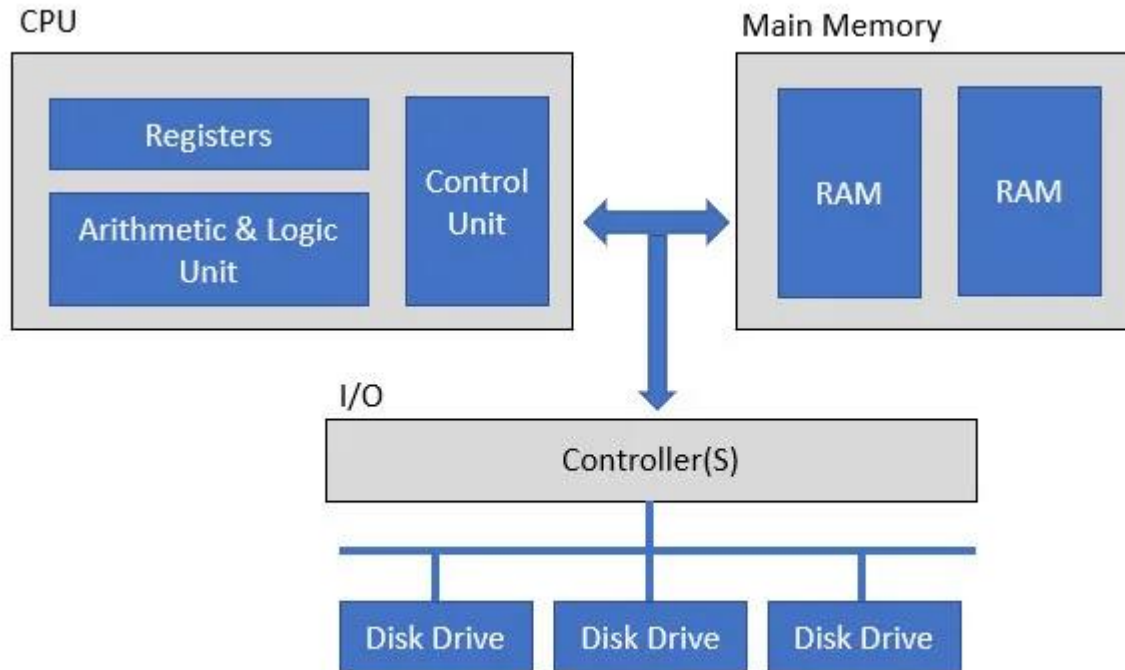
CALL (Compiler, Assembler, Linker, Loader)



CALL (Compiler, Assembler, Linker, Loader)

- Compiler: Converts C code into assembly code (RISC-V, x86)
- Assembler: Converts assembly code into machine code (raw bits)
 - Think 61C's RISC-V "green sheet"
- Linker: Deals with dependencies and libraries
 - You can ignore this part for 161
- Loader: Sets up memory space and jumps into the machine code
 - Sometimes, there is an additional linking step right before loading, called dynamic loading
 - Result: The executable doesn't need to contain all the dependencies
 - This also provides additional mitigations, as we'll see later
- After these steps, execution begins, and C runtime library calls **main!**

CPU, Memory, and Registers




Memory Layout

Textbook Chapter 2.3 & 2.5

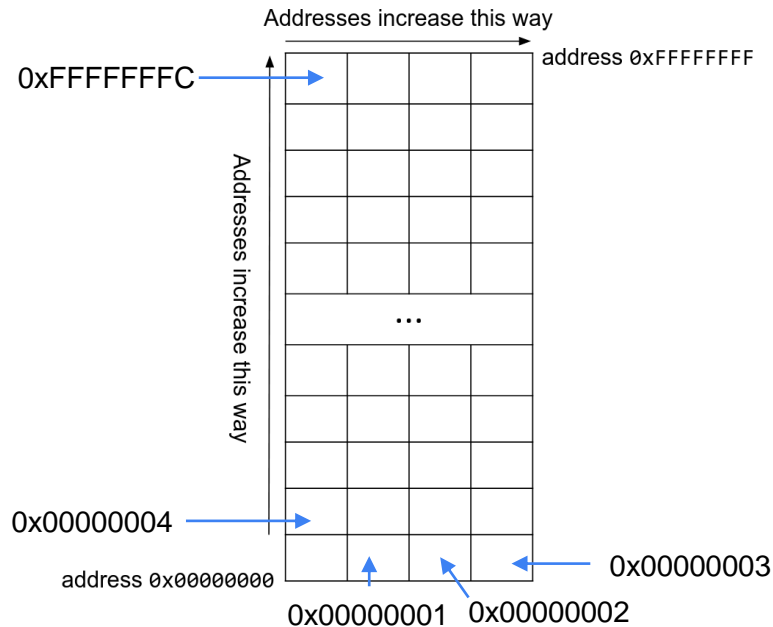
C Memory Layout

- At runtime, the loader tells your OS to give your program a big blob of memory
- On a 32-bit system, the memory has 32-bit addresses
 - On a 64-bit system, memory has 64-bit addresses
 - We use 32-bit systems in this class
- Each address refers to **one byte**, which means you have 2^{32} **bytes** of memory



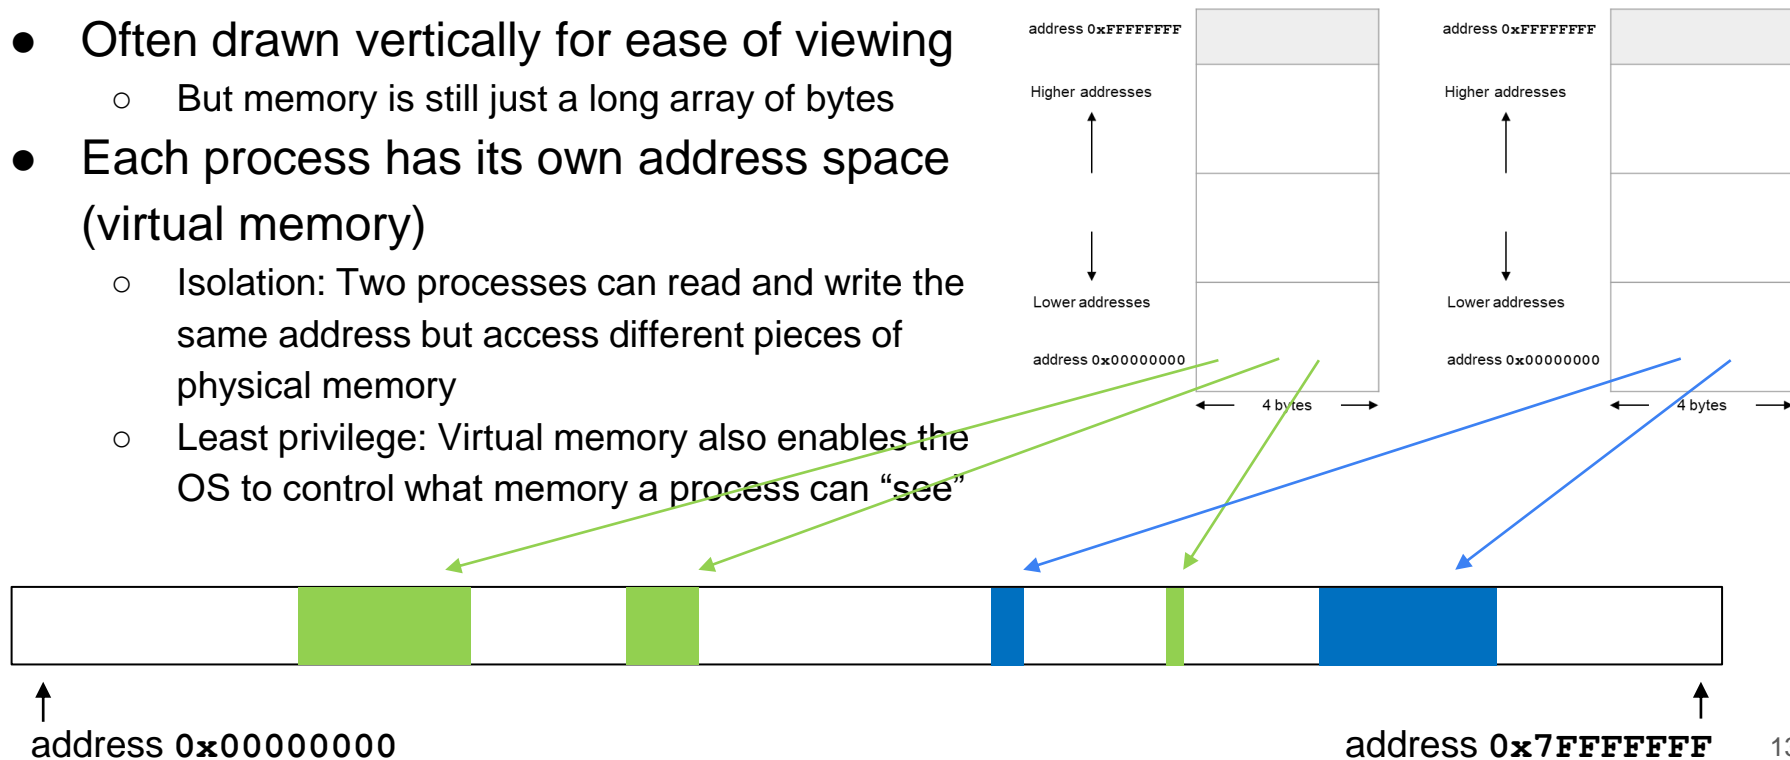
C Memory Layout

- Often drawn vertically for ease of viewing
 - But memory is still just a long array of bytes



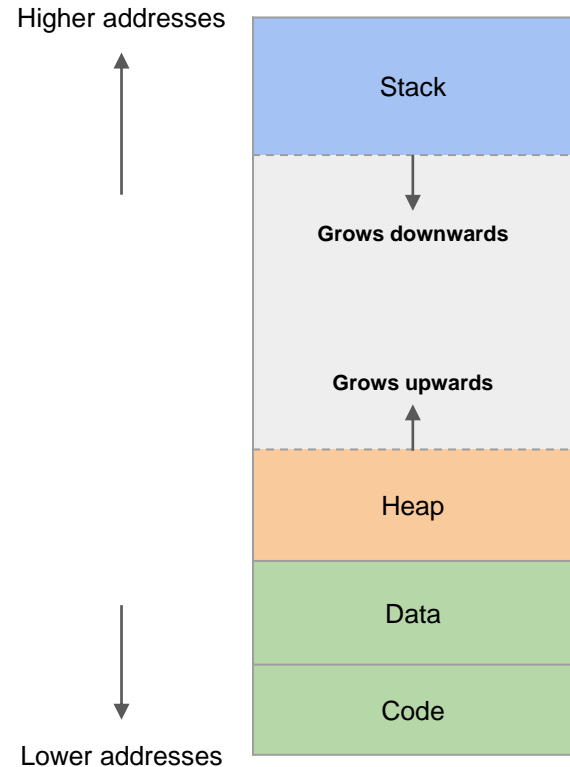
C Memory Layout

- Often drawn vertically for ease of viewing
 - But memory is still just a long array of bytes
- Each process has its own address space (virtual memory)
 - Isolation: Two processes can read and write the same address but access different pieces of physical memory
 - Least privilege: Virtual memory also enables the OS to control what memory a process can “see”



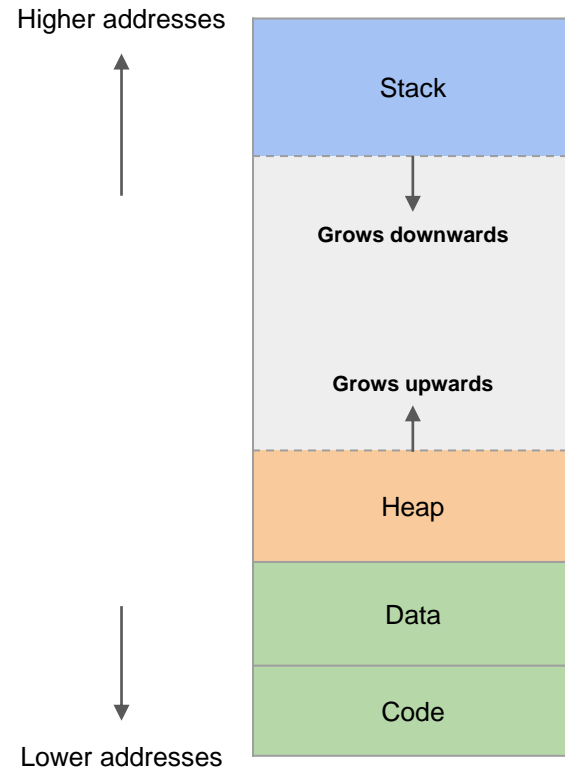
x86 Memory Layout

- Code
 - The program code itself (also called “text”)
- Data
 - Static variables, allocated when the program is started
- Heap
 - Dynamically allocated memory using **malloc** and **free**
 - As more and more memory is allocated, it grows *upwards*
- Stack:
 - Local variables and stack frames
 - As you make deeper and deeper function calls, it grows *downwards*



Registers

- Recall registers from CS 61C
 - Examples of RISC-V registers: **a0**, **t0**, **ra**, **sp**
- Registers are located on the CPU
 - This is different from the memory layout
 - Memory: addresses are 32-bit numbers
 - Registers: addresses are names (**ebp**, **esp**, **eip**)



x86 Architecture


Textbook Chapter 2.4 & 2.7

Why x86?

- It's the most commonly used instruction set architecture in consumer computers!
 - You are probably using an x86 computer right now... unless you're on a phone, tablet, or M1 Mac
- You only need enough to be able to read it and know what is going on
 - We will make comparisons to RISC-V, but it's okay if you haven't taken 61C and don't know RISC-V; you don't need to understand the comparisons to understand x86
- However, if you have a choice, choose ARM
 - Significantly higher performance for the same cost
 - 64b ARM has some unique security features we will discuss later

What is x86?

- Complex instruction set computer (CISC) architecture
 - The opposite of reduced instruction set computer (RISC) architecture
 - There are a lot of instructions... The full ISA manual is over **5,000 pages long!**
 - We will not be teaching the full instruction set (obviously), just enough of the calling convention to be able to do Project 1
- Launched in 1978 with the Intel 8086 processor and eventually took over much of computing
 - Over 40 years ago! The ISA is very bloated because of this...
- 64-bit variant x86-64 launched in 1999 (but we won't study it)




x86 Fact Sheet

- Little-endian

- The least-significant byte of multi-byte numbers is placed at the first/lowest memory address
- Same as RISC-V

storing the word 0x44332211 in memory:



`#include <stdint.h>`
`uint8_t` (1 byte)
`uint32_t` (4 bytes)

```
int main(void) {
    uint32_t num = 0xdeadbeef;

    // This prints "deadbeef".
    printf("%x", num);

    // This prints "ef be ad de".
    uint8_t *bytes = (uint8_t *) &num;
    for (size_t i = 0; i < 4; i++) {
        printf("%x ", bytes[i]);
    }
}
```

0xef	0xbe	0xad	0xde
------	------	------	------

x86 Registers

- Storage units as part of the CPU architecture (not part of memory)
- Only 8 main general-purpose registers:
 - EAX, EBX, ECX, EDX, ESI, EDI: General-purpose
 - ESP: Stack pointer (similar to `sp` in RISC-V)
 - EBP: Base pointer (similar to `fp` in RISC-V)
 - We will discuss ESP and EBP in more detail later
- Instruction pointer register: EIP
 - Similar to PC in RISC-V

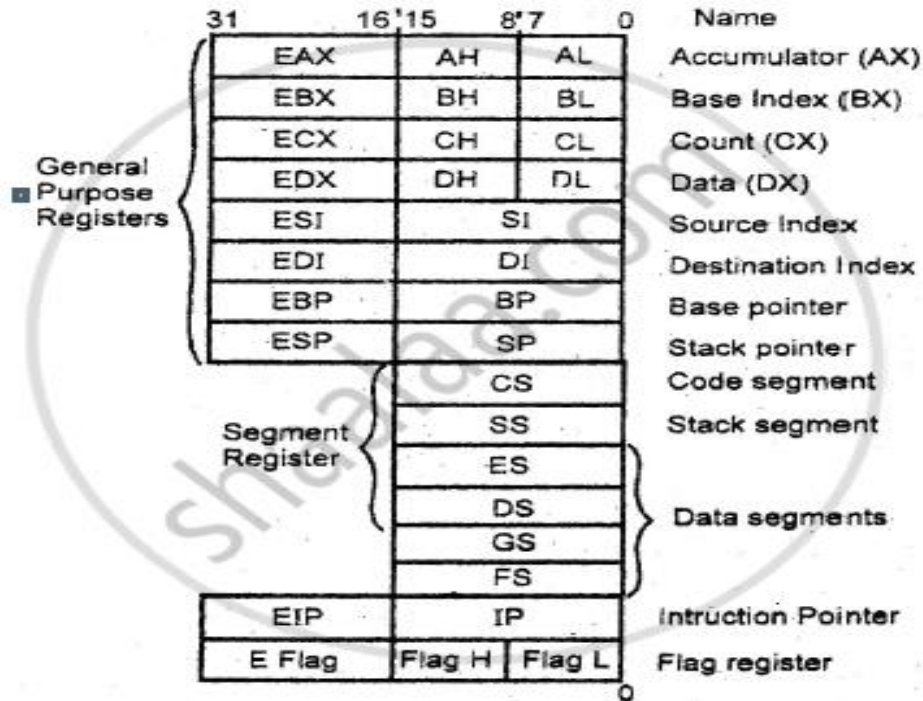
The main x86 registers...

- General purpose: EAX-EDX (E stands for Extended)
 - What you use for computing and other stuff, sorta...
- Indexes & Pointers
 - **EBP**: “Frame pointer”: points to the top/start of the current call frame on the stack
 - **ESP**: “Stack pointer”: points to the current stack
 - (Remember, stack grows down!)
 - PUSH and POP
 - Decrement the stack pointer and store something there
 - Load something and increment the stack pointer
 - Most operations are done with data on the stack...

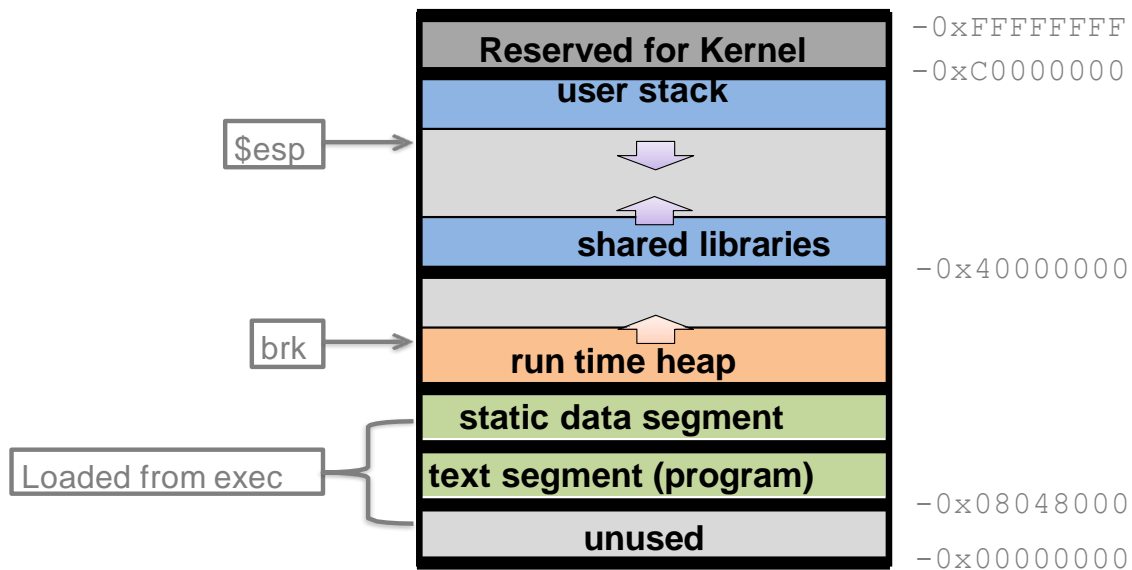
<https://www.cs.rutgers.edu/~pxk/419/notes/frames.html>

The main x86 registers...

- General purpose: EAX-EDX
 - What you use for computing and other stuff
- Indexes & Pointers
 - **EBP**: “Frame pointer”: points to the top/stack
 - **ESP**: “Stack pointer”: points to the current
 - (Remember, stack grows down!)
 - PUSH and POP
 - Decrement the stack pointer and store something there
 - Load something and increment the stack pointer
 - Most operations are done with data on the stack



Linux (32-bit) process memory layout




x86 Syntax

- Register references are preceded with a percent sign %
 - Example: `%eax`, `%esp`, `%edi`
- immediates are preceded with a dollar sign \$
 - Example: `$1`, `$161`, `$0x4`
- Memory references use parentheses and can have immediate offsets
 - Example: `8(%esp)` dereferences memory 8 bytes above the address contained in ESP

x86 Assembly

- Instructions are composed of an opcode and zero or more operands.


• **add \$0x8 , %ebx**



- Pseudocode: **EBX = EBX + 0x8**
- The destination comes last
 - Contrast with RISC-V assembly, where the destination (RD) is first
- The **add** instruction only has two operands; and the destination is an input
 - Contrast with RISC-V, where the two source operands are separate (RS1 and RS2)
- This instruction uses a register and an immediate

x86 Assembly

- **xor 4(%esi), %eax**



- Pseudocode: **EAX = EAX ^ *(ESI + 4)**
- This is a memory reference, where the value at 4 bytes above the address in ESI is dereferenced, XOR'd with EAX, and stored back into EAX
 - Most instructions can be register-register, register-immediate, register-memory, or memory-immediate (but not memory-memory)
 - How can you achieve a memory-memory operation?

Stack Layout


Textbook Chapter 2.6

Stack Frames

- When your code calls a function, space is made on the stack for local variables
 - This space is known as the **stack frame** for the function
 - The stack frame goes away once the function returns
- The stack starts at higher addresses. Every time your code calls a function, the stack makes extra space by growing down
 - Note: Data on the stack, such as a string, is still stored from lowest address to highest address. “Growing down” only happens when extra memory needs to be allocated.



Stack Frames

- To keep track of the current stack frame, we store two pointers in registers
 - The EBP (base pointer) register points to the top of the current stack frame
 - Equivalent to RISC-V **fp**
 - The ESP (stack pointer) register points to the bottom of the current stack frame
 - Equivalent to RISC-V **sp** (but x86 moves the stack pointer up and down a lot more than RISC-V does)





Pushing and Popping

- The **push** instruction adds an element to the stack
 - Decrement ESP to allocate more memory on the stack
 - Save the new value on the lowest value of the stack



Pushing and Popping

- The **pop** instruction removes an element from the stack
 - Load the value from the lowest value on the stack and store it in a register
 - Increment ESP to deallocate the memory on the stack



Why push and pop?



- This allows x86 to use “stack-machine”-style logic
 - Operations always operate on the top items on the stack and push the result back onto the stack
 - Example: $a + b * c + d$

■ `PUSH a` // `a`

■ `PUSH b` // `a, b`

■ `PUSH c` // `a, b, c`

■ `MUL` // `a, (b*c)`

■ `ADD` // `(a+(b*c))`

■ `PUSH` // `(a+(b*c)), d`

■ `ADD` // `(a+(b*c))+d`

■ Result is the last remaining item on the stack!

First convert the expression to postfix notation, then apply stack operations.

<https://www.geeksforgeeks.org/arithmetic-expression-evaluation/>

- Used to think this allowed for smaller code and simpler code generation
 - Reality? Not so much. RISC logic is equally efficient in code size in practice, especially with the compressed instructions for RISC-V.
 - But this misconception is why the Java VM is a stack machine too...
- (Not really a) Takeaway: Take 164 if you want to learn more about this

x86 Stack Layout

- Local variables are always allocated on the stack
 - Contrast with RISC-V, which has plenty of registers that can be used for variables
- Individual variables within a stack frame are stored with the first variable at the *highest* address
- Members of a struct are stored with the first member at the *lowest* address
- Global variables (not on the stack) are stored with the first variable at the *lowest* address


Stack Layout

```
struct foo {  
    long long f1; // 8 bytes  
    int f2;       // 4 bytes  
    int f3;       // 4 bytes  
};  
  
void func(void) {  
    int a;        // 4 bytes  
    struct foo b;  
    int c;        // 4 bytes  
}
```

Higher addresses



Lower addresses



← 4 bytes →

How would you fill out the boxes in this stack diagram?

Options:

a b.f1 b.f2 b.f3 c